Robust Loss Development Using MCMC: A Vignette

Christopher W. Laws   Frank A. Schmid

July 2, 2010

Abstract

For many lines of insurance, the ultimate loss associated with a particular exposure (accident or policy) year may not be realized (and hence known) for many calendar years; instead these losses develop as time progresses. The actuarial concept of loss development aims at estimating (at the level of the aggregate loss triangle) the ultimate loss by exposure year, given their respective stage of maturity (as defined by the time distance between the exposure year and the latest observed calendar year). This vignette describes and demonstrates loss development using of the package `lossDev`, which centers on a Bayesian time series model. Notable features of this model are a skewed Student-t distribution with time-varying scale and skewness parameters, the use of an expert prior for the calendar year effect, and the ability to accommodate a structural break in the consumption path of services. R and the package are open-source software projects and can be freely downloaded from CRAN: `http://cran.r-project.org` and `http://lossdev.r-forge.r-project.org/`.

1 Installation

At the time of writing this vignette, the current version of `lossDev` is 0.9.3, which has been released as an R package and can be downloaded from `http://lossdev.r-forge.r-project.org/`. `lossDev` should be available on CRAN shortly. (For instructions on installing R packages please see the help files for R.) `lossDev` requires `rjags` for installation. `rjags` requires that a valid version of JAGS be installed on the system. JAGS is an open source program for analysis of Bayesian hierarchical models using Markov Chain Monte Carlo (MCMC) simulation and can be freely download from `http://calvin.iarc.fr/~martyn/software/jags/`.

2 Model Overview

`lossDev` identifies three time dimensions in the data-generating process of the loss triangle. Specifically, the incremental payments are driven by three time series processes, which manifest themselves in exposure growth, development, and the calendar year effect; these processes are illustrated in Figure 1.

In the model, the growth rate that represents the calendar year effect is denoted κ. The rate of exposure growth, η, is net of the calendar year effect. The growth rate δ is the rate of decay in incremental payments, adjusted for the calendar year effect. Incremental payments that have been adjusted for the calendar year effect (and, hence, inflation) represent consumption of units of services; for instance, for an auto bodily injury triangle, this consumption pertains to medical services. A decline in consumption at the level of the aggregate loss triangle may be due to claimants exiting or due to remaining claimants decreasing their consumption.

For a more detailed explanation, including model equations, please see
Schmid, Frank A. "Robust Loss Development Using MCMC," 2009.

`lossDev` currently provides two models, both of which are designed to develop annual loss triangles.
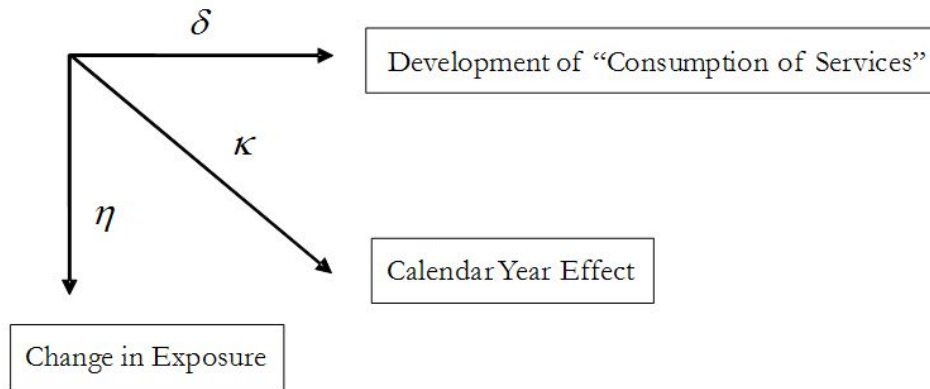
Figure 1: Triangle Dynamics.

Section 3 uses the first ("standard") model, which assumes that all exposure years are subject to a common consumption path.

Section 4 uses the second ("change point") model to develop a loss triangle with a structural break in the consumption path, thus assuming that earlier exposure years are subject to one consumption path and later exposure years are subject to another.

3   Using the Standard Model for Estimation

3.1   Data

The standard model (which does not allow for a structural break) is demonstrated on a loss triangle from Automatic Facilitative business in General Liability (excluding Asbestos & Environmental). The payments are on an incurred basis.

This triangle is taken from
Mack, Thomas, "Which Stochastic Model is Underlying the Chain Ladder Method," *Casualty Actuarial Society Forum*, Fall 1995, pp. 229-240, `http://www.casact.org/pubs/forum/95fforum/95ff229.pdf`.

3.2   Model Specification

Standard models are specified with the function `makeStandardAnnualInput`. This function takes as input all data used in the estimation process. `makeStandardAnnualInput` also allows the user to vary the model specification through several arguments. Most of these arguments have defaults that should be suitable for most purposes.

To ensure portability, the data used in this vignette is packaged in `lossDev` and as such is loaded using the `data` function. However, the user wishing to develop other loss triangles should load the data using standard R functions (such as `read.table` or `read.csv`). See the R manual for assistance.

3.2.1   Loading and Manipulating the Data

The Triangle   As input, `makeStandardAnnualInput` can take either a cumulative loss triangle or an incremental loss triangle (or in the case where one might not be directly calculable from the other, both triangles may be supplied). `makeStandardAnnualInput` expects any supplied loss triangle to be a matrix. The row names for the matrix must be the Accident (or Policy) Year

and must appear in ascending order. The matrix must be square and *all* values below the latest observed diagonal must be missing; missing values on and above this diagonal are permitted.

Note the negative value in Accident Year 1982 in the example triangle. Because incremental payments are modeled on the log scale, this value will be treated as missing, which could result in a slightly overstated ultimate loss. A comparison of predicted vs observed cumulative payments in Figure 9 indicates that, at least in this instance, this possible overstatement is not a concern.

```
> library(lossDev)

module basemod loaded
module bugs loaded
module lossDev loaded

> data(IncrementalGeneralLiablityTriangle)
> IncrementalGeneralLiablityTriangle <- as.matrix(IncrementalGeneralLiablityTriangle)
> print(IncrementalGeneralLiablityTriangle)

     DevYear1 DevYear2 DevYear3 DevYear4 DevYear5 DevYear6 DevYear7 DevYear8
1981     5012     3257     2638      898     1734     2642     1828      599
1982      106     4179     1111     5270     3116     1817     -103      673
1983     3410     5582     4881     2268     2594     3479      649      603
1984     5655     5900     4211     5500     2159     2658      984       NA
1985     1092     8473     6271     6333     3786      225       NA       NA
1986     1513     4932     5257     1233     2917       NA       NA       NA
1987      557     3463     6926     1368       NA       NA       NA       NA
1988     1351     5596     6165       NA       NA       NA       NA       NA
1989     3133     2262       NA       NA       NA       NA       NA       NA
1990     2063       NA       NA       NA       NA       NA       NA       NA
     DevYear9 DevYear10
1981       54       172
1982      535        NA
1983       NA        NA
1984       NA        NA
1985       NA        NA
1986       NA        NA
1987       NA        NA
1988       NA        NA
1989       NA        NA
1990       NA        NA
```

The Stochastic Inflation Prior Incremental payments may be subject to inflation. One can supply `makeStandardAnnualInput` with a price index, such as the CPI, as an expert prior for the rate of inflation. The supplied rate of inflation must cover the years of the supplied incremental triangle and may extend (both into the past and future) beyond these years. If a future year's rate of inflation is needed but is yet unobserved, it will be simulated from an Ornstein–Uhlenbeck process that has been calibrated to the supplied inflation series.

For this example, the CPI is taken as a prior for the stochastic rate of inflation.

Note that observed rates of inflation that extend beyond the last observed diagonal in `IncrementalGeneralLiablityTriangle` are not utilized in this example, although `lossDev` is capable of doing so.

```
> data(CPI)
> CPI <- as.matrix(CPI)[, 1]
```

```
> CPI.rate <- CPI[-1]/CPI[-length(CPI)] - 1
> CPI.rate.length <- length(CPI.rate)
> print(CPI.rate[(-10):0 + CPI.rate.length])

      1997       1998       1999       2000       2001       2002       2003
0.02294455 0.01557632 0.02208589 0.03361345 0.02845528 0.01581028 0.02279044
      2004       2005       2006       2007
0.02663043 0.03388036 0.03225806 0.02848214

> CPI.years <- as.integer(names(CPI.rate))
> years.available <- CPI.years <= max(as.integer(dimnames(IncrementalGeneralLiablityTriangle)[[1]
> CPI.rate <- CPI.rate[years.available]
> CPI.rate.length <- length(CPI.rate)
> print(CPI.rate[(-10):0 + CPI.rate.length])

      1980       1981       1982       1983       1984       1985       1986
0.13498623 0.10315534 0.06160616 0.03212435 0.04317269 0.03561116 0.01858736
      1987       1988       1989       1990
0.03649635 0.04137324 0.04818259 0.05403226
```

3.2.2 Selection of Model Options

The function `makeStandardAnnualInput` has many options to allow for customization of model specification; however, not all options will be illustrated in this tutorial.

For this example, the loss history is supplied as incremental payments to the argument `incremental.payments`. The exposure year type of this triangle is set to Accident Year by setting the value of `exp.year.type` to "ay." The default is "ambiguous" which should be sufficient in most cases as this information is only utilized by a handful of functions and the information can be supplied (or overridden calling those functions).

The function allows for the specification of two rates of inflation (in addition to a zero rate of inflation). One of these rates is allowed to be stochastic, meaning that uncertainty in future rates of this inflation series are simulated from a process calibrated to the observed series. For the current demonstration, it will be assumed that the CPI is the only applicable inflation rate, and that this rate is stochastic. This is done by setting the value of `stoch.inflation.rate` to `CPI.rate` (which was created earlier). The user has the option of specifying what percentage of dollars inflate at `stoch.inflation.rate`, with this value being allowed to vary for each cell of the triangle. For the current illustration, it is assumed that all dollars (in all periods) follow the CPI. This is done by setting `stoch.inflation.weight` to 1 and `non.stoch.inflation.weight` to 0.

By default, the measurement equation for the logarithm of the incremental payments is a Student-t. The user has the option of using a skewed-t by setting the value of `use.skew.t` to TRUE. For this demonstration, a skewed-t will be used.

Because `lossDev` is designed to develop loss triangles to ultimate, some assumptions must be made with regard to the extension of the consumption path beyond the number of development years in the observed triangle. The default assumes the last estimated decay rate (i.e., growth rate of consumption) is applicable for all future development years, and such is assumed for this example. This default can be overridden by the argument `projected.rate.of.decay`. Additionally, either the final number of (possibly) non-zero payments must be supplied via the argument `total.dev.years` or the number of non-zero payments in addition to the number of development years in the observed triangle must be supplied via the argument `extra.dev.years`. Similarly, the number of additional, projected exposure years can also be specified.

```
> standard.model.input <- makeStandardAnnualInput(incremental.payments = IncrementalGeneralLiabli
+     stoch.inflation.weight = 1, non.stoch.inflation.weight = 0,
```

```
+        stoch.inflation.rate = CPI.rate, exp.year.type = "ay", extra.dev.years = 5,
+        use.skew.t = TRUE)
```

3.3   Estimating the Model

Once the model has been specified, it can be estimated.

MCMC Overview The model is Bayesian and estimated by means of Markov chain Monte Carlo Simulation (MCMC). To perform MCMC, a Markov chain is constructed in such a way that the limiting distribution of the chain is the posterior distribution of interest. The chain is initialized with starting values and then run until it has reached a point of convergence in which samples adequately represent random (albeit sequentially dependent) draws from this posterior distribution. The set of iterations performed (and discarded) until samples are assumed to be draws from the posterior is called a "burn-in." After the burn-in, the chain is iterated further to collect samples. The samples are then used to calculated the statistic of interest.

While the user is not responsible for the construction of the Markov chain, he is responsible for assessing the chains' convergence. (Section 3.4.1 gives some pointers on this.) The most common way of accomplishing this task is to run several chains simultaneously with each chain having been started with a different set of initial values. Once all chains are producing similar results, one can assume that the chains have converged.

To estimate the model, the function `runLossDevModel` is called with the first argument being the input object created by `makeStandardAnnualInput`. To specify the number of iterations to discard, the user sets the value of `burnIn`. To specify the number of iterations to perform after the burn-in, set the value of `sampleSize`. To set the number of chains to run simultaneously, supply a value for `nChains`. The default value for `nChains` is 3, which should be sufficient for most cases.

It is also common practice (due to possible autocorrelation in the samples) to apply "thining," which means that only every n-th draw is stored. The argument `thin` is available for this purpose.

Memory Issues MCMC can require large amounts of memory. To allow `lossDev` to work with limited hardware, the R package `filehash` is used to cache the codas of monitored values to the hard-drive in an efficient way. While such caching can allow estimation of large triangles on computers with limited memory, it can also slow down some computations. The user has the option of turning this feature on and off. This is done via the function `lossDevOptions` by setting the argument `keepCodaOnDisk` to TRUE or FALSE.

R also makes available the function `memory.limit`, which one may find useful.

```
> standard.model.output <- runLossDevModel(standard.model.input,
+        burnIn = 30000, sampleSize = 30000, thin = 10)


Compiling data graph
   Resolving undeclared variables
   Allocating nodes
   Initializing
   Reading data back into data table
Compiling model graph
   Resolving undeclared variables
   Allocating nodes
   Graph Size: 7388

[1] "Update took 16.7776 mins"
```

3.4 Examining Output

`makeStandardAnnualInput` returns a complex output object. `lossDev` provides several user-level functions to access the information contained in this object. Many of these functions are described below.

3.4.1 Assessing Convergence

As mentioned, the user is responsible for assessing the convergence of the Markov chains used to estimate the model. To this aim, `lossDev` provides several functions to produce trace and density plots.

Arguably, the most important charts for assessing convergence are the trace plots associated with the three time dimensions of the model. Convergence of exposure growth, the consumption path, and the calendar year effect are assessed in Figures 2, 3, and 4 respectively. These charts are produced with the functions `exposureGrowthTracePlot`, `consumptionPathTracePlot`, and `calendarYearEffectErrorTracePlot`.

```
> exposureGrowthTracePlot(standard.model.output)
```
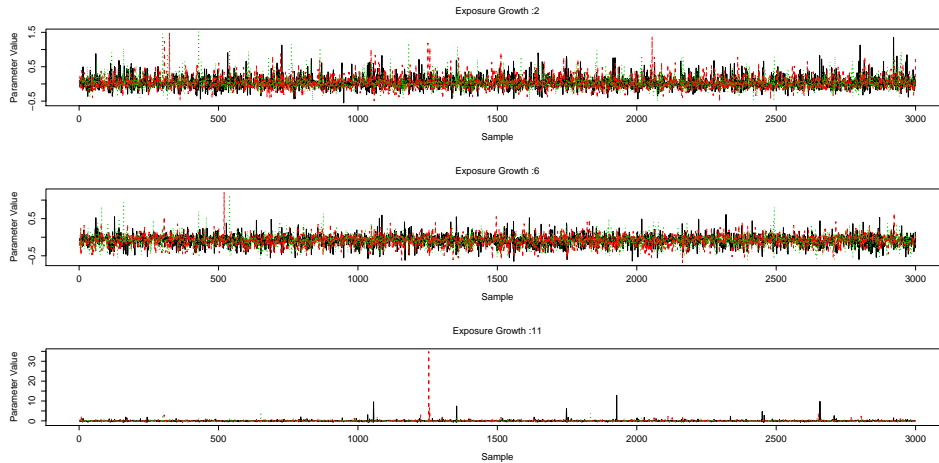

Figure 2: Trace plots for select exposure growth parameters.

3.4.2 Assessing Model Fit

`lossDev` provides many diagnostic charts to asses how well the model fits the observed triangle.

Residuals For the analysis of residuals, `lossDev` provides the function `triResi`. `triResi` plots the residuals (on the log scale) by the three time dimensions. The time dimension is selected by means of the argument `timeAxis`. By default, residual charts are standardized to account for any assumed/estimated heteroskedasticity in the (log) incremental payments. These charts can be found in Figures 5, 6, and 7.

Note that because (the log) incremental payments are allowed to be skewed, the residuals need not be symmetric.

QQ-Plot `lossDev` provides a QQ-Plot in the function `QQPlot`. `QQPlot` plots the median of simulated incremental payments (sorted at each simulation) against the observed incremental payments. Plotted points from a well calibrated model will be close to the 45-degree line. These results are shown in Figure 8.

```
> consumptionPathTracePlot(standard.model.output)
```
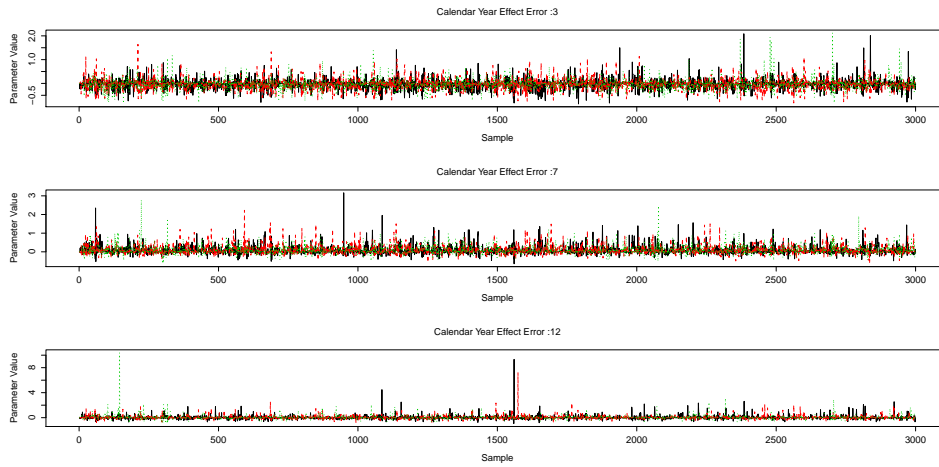

Figure 3: Trace plots for select development years on the consumption path.

```
> calendarYearEffectErrorTracePlot(standard.model.output)
```
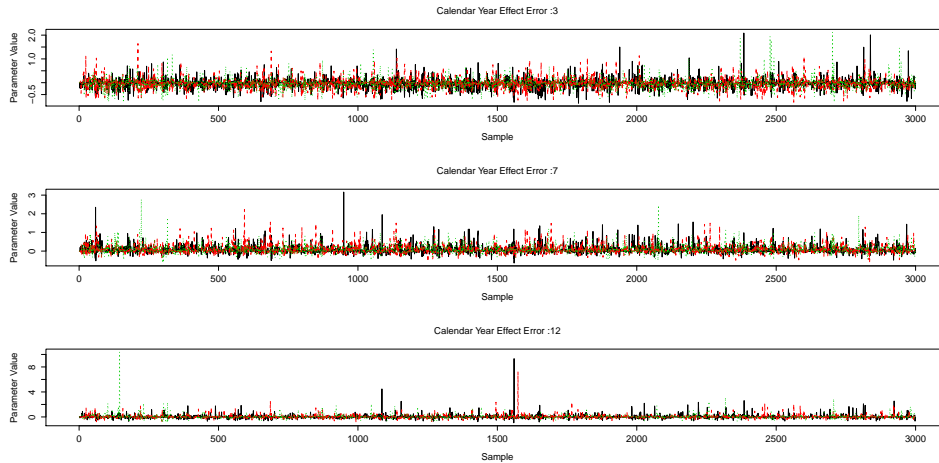

Figure 4: Trace plots for select calendar year effect errors.

Comparison of Cumulative Payments   As a means of assessing how well the predicted cumulative payments line up with the observed values, `lossDev` provides the function `finalCumulativeDiff`. This function plots the relative difference between the predicted and observed cumulative payments (when such payments exists) for the last observed cumulative payment in each exposure year, alongside credible intervals. These relative differences, which are shown in Figure 9, can be useful for assessing the impact of negative incremental payments, as discussed.

3.4.3   Extracting Inference and Results

After compiling, burning-in, and sampling, the user will wish to extract results from the output. Many of the functions mentioned in this section also return the values of some plotted information. These values are returned invisibly and as such are not printed at the REPL unless such an operation is requested. Additionally, many of these functions also provide an option to suppress plotting.

```
> triResi(standard.model.output, timeAxis = "dy")
```


Figure 5: Residuals by development year.

```
> triResi(standard.model.output, timeAxis = "ey")
```


Figure 6: Residuals by exposure year.

Predicted Payments   Perhaps the most practically useful function is `predictedPayments`. This function can plot and return the estimated incremental predicated payments. As the function can also plot the observed values against the predicted values (`plotObservedValues`), it also serves as a diagnostic tool. The log incremental payments are plotted against the predicted values in Figure 10.

`predictedPayments` can also plot and return the estimated cumulative payments and has the option of taking observed payments at "face value" (meaning that predicted payments are replaced with observed payments whenever possible) in the returned calculations; this can be useful for the construction of reserve estimates. In Figure 11, only the predicted cumulative payments are plotted. The function is also used to construct an estimate (with credible intervals) of the ultimate loss.

```
> standard.ult <- predictedPayments(standard.model.output, type = "cumulative",
+     plotObservedValues = FALSE, mergePredictedWithObserved = TRUE,
```

8

```
> triResi(standard.model.output, timeAxis = "cy")
```


Figure 7: Residuals by calendar year.

```
> QQPlot(standard.model.output)
```


Figure 8: QQ-Plot.

```
+       logScale = TRUE, quantiles = c(0.025, 0.5, 0.0975), plot = FALSE)
> standard.ult <- standard.ult[, , dim(standard.ult)[3]]
> print(standard.ult)
```

	1981	1982	1983	1984	1985	1986	1987	1988
2.5%	18863.61	16765.05	23645.15	27509.62	27057.34	17243.95	14499.92	16321.23
50%	19132.81	17226.12	24409.09	28835.41	29323.53	20780.52	19765.60	24016.68
9.75%	18913.01	16859.42	23822.13	27825.60	27676.75	18219.33	15955.44	18487.87

	1989	1990	1991
2.5%	9508.532	6447.192	3895.809
50%	19921.097	20258.532	21466.862
9.75%	12254.688	9682.883	7668.809

9

```
> finalCumulativeDiff(standard.model.output)
```


Figure 9: Difference in Final Observed Cumulative Payments.

```
> predictedPayments(standard.model.output, type = "incremental",
+       logScale = TRUE)
```


Figure 10: Predicted Incremental Payments.

Consumption Path `lossDev` makes the consumption path available via `consumptionPath`. The consumption path is the trajectory of exposure-adjusted and calendar year effect-adjusted log incremental payments and is modeled as a linear spline. The standard model assumes a common consumption path for all exposure years in the triangle. The use of this function is demonstrated in Figure 12; the displayed consumption path represents the exposure level of the first exposure year in the triangle.

Knots in the Consumption Path The consumption path is modeled as a linear spline. The number of knots in this spline is endogenous to the model. The function `numberOfKnots` can be used to extract information regarding the posterior number of knots. All else equal, a higher number of knots indicates a higher degree of non-linearity. Figure 13 illustrates the use of this function.

```
> predictedPayments(standard.model.output, type = "cumulative",
+     plotObservedValues = FALSE, logScale = TRUE)
```


Figure 11: Predicted Cumulative Payments.

```
> consumptionPath(standard.model.output)
```


Figure 12: Consumption Path.

Rate of Decay While the consumption path illustrates the level of exposure-adjusted and calendar year effect-adjusted log incremental payments, sometimes one may prefer to examine the development time force in terms of a decay rate. The rate of decay from one development year to the next (which is approximately the slope of the consumption path) is made available via the function `rateOfDecay`. As the standard model assumes a common consumption path for all exposure years, the standard model has only a single decay rate vector. An example of this function can be found in Figure 14.

Exposure Growth The year over year changes in the estimated exposure level are made available by the function `exposureGrowth`. An example of this function can be found in Figure 15.

Calendar Year Effect The model assumes that the cells on a diagonal are subject to a correlated shock. The shock consists of a component exogenous to the triangle (generally represented

11

```
> numberOfKnots(standard.model.output)
```


Figure 13: Number of Knots.

```
> rateOfDecay(standard.model.output)
```


Figure 14: Rate Of Decay.

by a price index, such as the CPI) and an endogenous stochastic component. This endogenous component is the calendar year effect error, defined as the difference between the estimated calendar year effect and the expert prior for the rate of inflation. As `lossDev` allows the user to vary the exogenous component for each cell, graphically displaying the entire calendar year effect requires three dimensions. This is done by plotting a grid of colored blocks and varying the intensity of each color according to the associated calendar year effect. An example of this can be found in Figure 16. Note that the value in cell (1,1) is undefined.

Alternatively, one could merely plot the endogenous stochastic component. As this calendar year effect error is common to all cells on a given diagonal, the number of dimensions is reduced by one. An illustration of the calendar year effect error is displayed in Figure 17. In this example, the calendar year effect error displays a fair degree of autocorrelation. `lossDev` can account for such correlation by setting the argument `use.ar1.in.calendar.year` in `makeStandardAnnualInput` to TRUE. Exploring this is left as an exercise to the reader.

```
> exposureGrowth(standard.model.output)
```


Figure 15: Exposure Growth.

```
> calendarYearEffect(standard.model.output)
```



Figure 16: Calendar Year Effect.

**Changes In Variance**   As development time progresses, the number of transactions that comprise a given incremental payment declines. This can lead to an increase in the variance of the log incremental payments even as the level of the payments may decrease. In order to account for this potential increase in variance, the model (optionally) allows for the scale parameter of the Student-$t$ to vary with development time. This scale parameter is smoothed via a second-order random walk on the log scale. As a result, the standard deviation can vary for each development year. An example is displayed in Figure 18.

**Skewness Parameter**   The measurement equation for the log incremental payments is (optionally) a skewed-$t$. skewnessParameter allows for the illustration of the posterior skewness parameter. (For reference, the prior is also illustrated.) While the skewness parameter does not directly translate into the estimated skewness, the two are related. For instance, a skewness parameter of zero would correspond to zero skew. An example is displayed in Figure 19.

```
> calendarYearEffectErrors(standard.model.output)
```



Figure 17: Calendar Year Effect Errors.

```
> standardDeviationVsDevelopmentTime(standard.model.output)
```



Figure 18: Standard Deviation vs Development Time.

**Degrees of Freedom**  The degrees of freedom associated with the measurement equation is endogenous to the model estimation. To ensure existence of moments, when estimating a skewed-$t$, the degrees of freedom is constrained to be greater than 4; otherwise this value is constrained to be greater than 2. All else equal, lower degrees of freedom indicate the presence of heavy tails.

The `lossDev` function `degreesOfFreedom` allows for the illustration of the posterior degrees of freedom. (For reference, the prior is also illustrated.) Figure 19 displays the posterior degrees of freedom for this example.

### 3.4.4  The Ornstein–Uhlenbeck Process

Future values for the assumed stochastic rate of inflation are simulated from an Ornstein–Uhlenbeck process. `lossDev` allows the user to examine predicted and forecast values as well as some of the underlying parameters. Such options are outlined below.

```
> skewnessParameter(standard.model.output)
```



Figure 19: Skewness Parameter.

```
> degreesOfFreedom(standard.model.output)
```



Figure 20: Degrees Of Freedom.

**Fit and Forecast**    To display the fitted values vs the observed values (as well as the forecast values) the user must use the function `stochasticInflation`. The chart for the example illustrated above is displayed in Figure 21.

**Stationary Mean**    The Ornstein–Uhlenbeck process has a stationary mean; disturbances from this mean are assumed to be correlated. Specifically, the projected rate of inflation will (geometrically) approach the stationay mean as time progresses. This stationary mean can be graphed with the function `StochasticInflationStationaryMean`. The chart for the example illustrated above is displayed in Figure 22.

**Autocorrelation**    The Ornstein – Uhlenbeck process assumes that the influence of a disturbance decays geometrically with time. The parameter governing this rate is traditionally referred to as $\rho$. To obtain this value, call the function `StochasticInflationRhoParameter`. The chart for the example illustrated above is displayed in Figure 23.

15

```
> stochasticInflation(standard.model.output)
```



Figure 21: Stochastic Inflation Fit.

```
> stochasticInflationStationaryMean(standard.model.output)
```



Figure 22: Estimated Stochastic Inflation Stationary Mean.

```
          used (Mb) gc trigger  (Mb) max used  (Mb)
Ncells  369772  9.9      667722  17.9     667722  17.9
Vcells 6295676 48.1    14762915 112.7 14762914 112.7

          used (Mb) gc trigger  (Mb) max used  (Mb)
Ncells  369765  9.9      667722  17.9     667722  17.9
Vcells 6295673 48.1    14762915 112.7 14762914 112.7
```

4   Using the Change Point Model for Estimation

The standard model outlined in Section 3 assumes the same consumption path for all exposure years. Due to changes in the loss environment, this may not be appropriate for all loss triangles. A triangle that may have experienced a structural break in the consumption path is outlined below.

```
> stochasticInflationRhoParameter(standard.model.output)
```


Figure 23: Estimated Stochastic Inflation Rho Parameter.

4.1 Data

The triangle used for this example is a Private Passenger Auto Bodily Injury Liability triangle and consists of accident year data on a paid basis.

In December 1986, a judicial decision limited the ability of judges to dismiss cases. This judicial decision may have brought about a change in the consumption path, thus making this triangle a good example for the change point model.

This triangle is taken from
Hayne, Roger M., "Measurement of Reserve Variability," *Casualty Actuarial Society Forum*, Fall 2003, pp. 141-172, http://www.casact.org/pubs/forum/03fforum/03ff141.pdf.

4.2 Model Specification

4.2.1 Loading and Manipulating the Data

The Triangle Section 3.2.1 supplied incremental payments as model input. For variety, cumulative payments are supplied in this example.

Note the large number of payments at zero amounts. Because the model will treat these payments as missing values (since they are equal to negative infinity on the log scale), the predicted payments may be overstated. This issue is addressed in Section 5.

```
> data(CumulativeAutoBodilyInjuryTriangle)
> CumulativeAutoBodilyInjuryTriangle <- as.matrix(CumulativeAutoBodilyInjuryTriangle)
> sample.col <- (dim(CumulativeAutoBodilyInjuryTriangle)[2] - 6:0)
> print(decumulate(CumulativeAutoBodilyInjuryTriangle)[1:7, sample.col])

     DevYear12 DevYear13 DevYear14 DevYear15 DevYear16 DevYear17 DevYear18
1974        20        25         0        12         0         0         0
1975        18        67         0         0         0        32        NA
1976        96         7         8        18         0        NA        NA
1977         2         0        50         0        NA        NA        NA
1978       -55        18        21        NA        NA        NA        NA
1979        19        26        NA        NA        NA        NA        NA
1980         5        NA        NA        NA        NA        NA        NA
```

17

The Stochastic Inflation Expert Prior The MCPI (Medical Care Component of the CPI) is chosen as a an expert prior for the stochastic rate of inflation. While in Section 3.2.1 the expert prior did not extend beyond the observed diagonals (for realism), here a few extra observed years of the MCPI inflation are used for illustration purposes.

```
> data(MCPI)
> MCPI <- as.matrix(MCPI)[, 1]
> MCPI.rate <- MCPI[-1]/MCPI[-length(MCPI)] - 1
> print(MCPI.rate[(-10):0 + length(MCPI.rate)])

        1997       1998       1999       2000       2001       2002       2003
0.02804557 0.03196931 0.03510946 0.04070231 0.04601227 0.04692082 0.04026611
        2004       2005       2006       2007
0.04375631 0.04224444 0.04022277 0.04418203

> MCPI.years <- as.integer(names(MCPI.rate))
> max.exp.year <- max(as.integer(dimnames(CumulativeAutoBodilyInjuryTriangle)[[1]]))
> years.to.keep <- MCPI.years <= max.exp.year + 3
> MCPI.rate <- MCPI.rate[years.to.keep]
```

4.2.2 Selection of Model Options

While `makeStandardAnnualInput` (Section 3.2.2) is used to specify models without a change point (i.e., structural break), `makeBreakAnnualInput` is used to specify models with a change point. `makeBreakAnnualInput` has most of its arguments in common with `makeStandardAnnualInput`, and all these common arguments carry their meanings forward. However, `makeBreakAnnualInput` adds a few new arguments; these are for specifying the location of the structural break.

Most notable is the argument `first.year.in.new.regime` which, as the name suggests, indicates the first year in which the new consumption path applies. This argument can be supplied with a single value, in which case the model will give a hundred percent probability that this year is the first year in the new regime. However, this argument can also be supplied with a range of contiguous years, and the model will then estimate the first year in the new regime. Because the possible break occurs in late 1986, the range of years chosen for this example is 1986 to 1987.

The prior for the first year in the new regime is a discretized beta distribution. The user has the option of choosing the parameters for this prior by setting the argument `prior.for.first.year.in.new.regime`. Here, since the change was in late 1986, we choose a prior that accords more probability to the later year.

The argument `bound.for.skewness.parameter` is set to 5. This avoids the MCMC chain from "getting stuck" in the lower tail of the distribution (in this particular example). One should use the function `skewnessParemeter` (Figure 37) to evaluate the need to set this value. If the user is experiencing difficulties with the skewed-t, he may wish to use the non-skewed-t by setting the argument `use.skew.t` equal to FALSE (which is the default).

```
> break.model.input <- makeBreakAnnualInput(cumulative.payments = CumulativeAutoBodilyInjuryTrian
+     stoch.inflation.weight = 1, non.stoch.inflation.weight = 0,
+     stoch.inflation.rate = MCPI.rate, first.year.in.new.regime = c(1986,
+         1987), prior.for.first.year.in.new.regime = c(2, 1),
+     exp.year.type = "ay", extra.dev.years = 5, use.skew.t = TRUE,
+     bound.for.skewness.parameter = 5)
```

4.3 Estimating the Model

Just like in Section 3.3, the S4 object returned by `makeBreakAnnualInput` must be supplied to the function `runLossDevModel` in order to produce estimates.

```
> break.model.output <- runLossDevModel(break.model.input, burnIn = 30000,
+     sampleSize = 30000, thin = 10)

Compiling data graph
   Resolving undeclared variables
   Allocating nodes
   Initializing
   Reading data back into data table
Compiling model graph
   Resolving undeclared variables
   Allocating nodes
   Graph Size: 19945

[1] "Update took 1.280564 hours"
```

4.4 Examining Output

4.4.1 Assessing Convergence

As discussed, the user must examine the MCMC runs for convergence using the same functions mentioned in Section 3.4.1. To avoid repetition, only a few of the previously illustrated charts will be discussed below.

Because the change point model has two consumption paths, the method `consumptionPathTracePlot` for output related to this model has an additional argument when it comes to specifing the consumption path. If the argument `preBreak` equals TRUE, then the trace for the consumption path relevant to exposure years prior to the structural break will be plotted. Otherwise, the trace for the consumption path relevant to exposure years after the break will be plotted.

The trace for the pre-break consumption path is plotted in Figure 24. The trace for the post-break path is plotted in Figure 25.

```
> consumptionPathTracePlot(break.model.output, preBreak = TRUE)
```


Figure 24: Trace plots for select development years on the pre-break consumption path.

4.4.2 Assessing Model Fit

All of the functions mentioned in Section 3.4.2 are available for the change point model as well.

19

```
> consumptionPathTracePlot(break.model.output, preBreak = FALSE)
```


Figure 25: Trace plots for select development years on the post-break consumption path.

Residuals   One feature of `triResi` not mentioned in Section 3.4.2 is the option to turn off the standardization. As discussed, the model accounts for an increase in the variance of incremental payments as development time progresses by allowing a scale parameter to vary with development time. By default, `triResi` accounts for this by standardizing all the residuals to have a standard deviation of one. Turning off this feature (via the argument `standardize`) can provide insight into this process.

The standardized residuals for the change point model are displayed by development time in Figure 26. Figure 27 shows the residuals without this standardization.

```
> triResi(break.model.output, timeAxis = "dy")
```


Figure 26: (Standardized) Residuals by development year.

Comparison of Cumulative Payments   As mentioned, the loss triangle used to illustrate the change point model has a non-negligible number of incremental payments at the zero amount. Figure 28 uses the function `finalCumulativeDiff` to examine the impact of treating these values as missing.

```
> triResi(break.model.output, standardize = FALSE, timeAxis = "dy")
```


Figure 27: (Unstandardized) Residuals by development year.

```
> finalCumulativeDiff(break.model.output)
```


Figure 28: Difference in Final Observed Cumulative Payments.

4.4.3 Extracting Inference and Results

As was done for the standard model, the user will want to draw inferences from the change point model. All of the functions discussed in Section 3.4.3 are available for this purpose–though some will plot slightly different charts and return answers in slightly different ways. In addition, a few functions are made available to deal with the change point. These functions have no meaning for the standard model discussed in Section 3.

Predicted Payments   Figure 29 again uses the function `predictedPayments` to plot the predicted incremental payments vs the observed incremental payments. The impact of treating incremental payments of zero as missing values is most noticeable in this chart.

Consumption Path   Figure 30 plots the consumption path for the change point model, again using the function `consumptionPath`. Note that now two consumption paths are plotted – one for

```
> predictedPayments(break.model.output, type = "incremental", logScale = TRUE)
```


Figure 29: Predicted Incremental Payments.

the pre-break path and one for the post-break path. Both the pre- and post- break paths represent the exposure level of the first exposure year.

```
> consumptionPath(break.model.output)
```


Figure 30: Consumption Path.

Knots in the Consumption Path   Figure 31 displays the posterior number of knots for the change point model example, again using the function `numberOfKnots`. Note that the number of knots of both the pre-break and the post-break consumption paths are plotted.

Rate of Decay   Figure 32 uses the function `rateOfDecay` to plot the rate of decay from one development year to the next for both the pre- and post- break regimes. This can be useful in assessing the impact of a structural break in the run-off.

Calendar Year Effect   Figure 33 uses the function `calendarYearEffect` to plot the calendar year effect for the change point model. By default, `calendarYearEffect` will plot the calendar year

```
> numberOfKnots(break.model.output)
```

Figure 31: Number of Knots.

```
> rateOfDecay(break.model.output)
```

Figure 32: Rate Of Decay.

effect for all (observed and projected) incremental payments. Setting the argument `restricted-Size` to TRUE will plot the calendar year effect for only the observed incremental payments and the projected incremental payments needed to "square" the triangle. This feature can be useful for insurance lines with long tails.

Figure 34 shows the calendar year effect error which is plotted using the function `calendarYear-EffectErrors`.

Autocorrelation in Calendar Year Effect   The autocorrelation exhibited in Figure 34 is too strong to ignore. Figure 35 illustrates the use of `makeBreakAnnualInput`'s argument `use.ar1.in.calendar.year`.

Setting `use.ar1.in.calendar.year` to TRUE enables the use of an additional function: `calendarYearEffectAutoregressiveParameter`. This function will plot the autoregressive parameter associated with the calendar year effect error. Figure 36 illustrates the use of this function.

```
> calendarYearEffect(break.model.output)
```


Figure 33: Calendar Year Effect.

```
> calendarYearEffectErrors(break.model.output)
```


Figure 34: Calendar Year Effect Errors (Without AR1).

Skewness Parameter   Figure 37 displays the skewness parameter for the change point model example by using the function `skewnessParemeter`. The result of setting `bound.for.skewness.parameter` to 5 is visible in the chart.

First Year in New Regime   The posterior for the first year in which the post-break consumption path applies can be obtained via the function `firstYearInNewRegime`. Figure 38 shows the posterior (and prior) for the first year in the new regime. Note how the choice of the argument `prior.for.first.year.in.new.regime` to `makeBreakAnnualInput` has affected the prior.

5   Accounting for Incremental Payments of Zero

As mentioned in Section 4.2.1 and illustrated in Figure 29, the triangle used as an example for the change point model contains several incremental payments of zero which, if ignored, could cause

```
> break.model.input.w.ar1 <- makeBreakAnnualInput(cumulative.payments = CumulativeAutoBodilyInjur
+     stoch.inflation.weight = 1, non.stoch.inflation.weight = 0,
+     stoch.inflation.rate = MCPI.rate, first.year.in.new.regime = c(1986,
+         1987), prior.for.first.year.in.new.regime = c(2, 1),
+     exp.year.type = "ay", extra.dev.years = 5, use.skew.t = TRUE,
+     bound.for.skewness.parameter = 5, use.ar1.in.calendar.year = TRUE)
> break.model.output.w.ar1 <- runLossDevModel(break.model.input.w.ar1,
+     burnIn = 30000, sampleSize = 30000, thin = 10)

Compiling data graph
    Resolving undeclared variables
    Allocating nodes
    Initializing
    Reading data back into data table
Compiling model graph
    Resolving undeclared variables
    Allocating nodes
    Graph Size: 26907

[1] "Update took 1.313121 hours"

> calendarYearEffectErrors(break.model.output.w.ar1)
```


Figure 35: Calendar Year Effect Errors (With AR1).

the predicted losses to be overestimated.

lossDev provides a means to account for these payments at the zero amount. This is done by estimating a secondary, auxiliary model to determine the probably that a payment will be greater than zero. Predicted payments are then weighted by this probability.

5.1   Estimating the Auxiliary Model

To account for payments at zero amounts, the function accountForZeroPayments is called with the first argument being an object returned from a call to runLossDevModel. This function will then return another object which, when called by certain functions already mentioned, will incorporate into the calculation the probability that any particular payment is zero.

```
> break.model.output.w.zeros <- accountForZeroPayments(break.model.output)
```

```
> calendarYearEffectAutoregressiveParameter(break.model.output.w.ar1)
```


Figure 36: Calendar Year Effect Autoregressive Parameter.

```
> skewnessParameter(break.model.output)
```


Figure 37: Skewness Parameter.

```
Compiling model graph
   Resolving undeclared variables
   Allocating nodes
   Graph Size: 3157

[1] "Update took 26.031 secs"
```

5.2 Assessing Convergence of the Auxiliary Model

The MCMC run used to estimate the auxiliary model must be checked for convergence. `lossDev` provides the function `gompertzParameters` to this end.

The auxiliary model uses a (two-parameter) gompertz function to model the incremental payments at the zero amount. Which of these parameters is plotted by `gompertzParameters` is determined by the argument `parameter`.

```
> firstYearInNewRegime(break.model.output)
```


Figure 38: First Year in New Regime.

Figure 39 plots the parameter that determines the steepness of the curve. This parameter can be examined by setting `parameter` equal to "scale."

```
> gompertzParameters(break.model.output.w.zeros, parameter = "scale")
```


Figure 39: Gompertz Scale Parameter.

Figure 40 plots the parameter that determines the point in development time at which the curve assigns equal probability to payments being zero and payments being greater than zero; this parameter can be examined by setting `parameter` equal to "fifty.fifty."

5.3 Assessing Fit of the Auxiliary Model

One can plot the observed empirical probabilities of payments being greater than zero against the predicted (and projected) probabilities. This is done with the function `probablityOfPayment`. Figure 41 plot this chart.

```
> gompertzParameters(break.model.output.w.zeros, parameter = "fifty.fifty")
```


Figure 40: Gompertz Location Parameter.

```
> probablityOfPayment(break.model.output.w.zeros)
```


Figure 41: Probability of Non-Zero Payment.

5.4 Incorporating the Probability of Non-Zero Payment

Once the auxiliary model has been estimated and its output verified, the functions `predicted-Payments`, `finalCumulativeDiff`, and `tailFactor` will incorporate this information into their calculations.

Figure 42 displays the predicted incremental payments after accounting for the probability that some of them may be zero. This should be compared with Figure 29, which does not account for the possibility that payments may be zero.

```
> predictedPayments(break.model.output.w.zeros, type = "incremental",
+     logScale = TRUE)
```


Figure 42: Predicted Incremental Payments (Accounting for Zero Payments).